The Dial-A-Ride Problem with Transfers

Renaud Masson, Fabien Lehuédé, Olivier Péton *

LUNAM Université, Institut de Recherche en Communications et Cybernétique de Nantes, Ecole des Mines de Nantes, IRCCyN UMR CNRS 6597, 4 rue Alfred Kastler, 44307 Nantes, France

ARTICLE INFO

ABSTRACT

The Dial-A-Ride Problem with Transfers (DARPT) consists in defining a set of routes that satisfy transportation requests of users between a set of pickup points and a set of delivery points, in the presence of ride time constraints. Users may change vehicles during their trip. This change of vehicle, called a transfer, is made at specific locations called transfer points. Solving the DARPT involves modeling and algorithmic difficulties. In this paper we provide a solution method based on an Adaptive Large Neighborhood Search (ALNS) metaheuristic and explain how to check the feasibility of a request insertion. The method is evaluated on real-life and generated instances. Experiments show that savings due to transfers can be up to 8% on real-life instances.

1. Introduction

The Dial-A-Ride Problem (DARP) consists in defining a set of minimum cost routes in order to satisfy a set of transportation requests. Each request involves transporting a set of users from a set of origins, called *pickup points* to a set of destinations, called *delivery points*. Users associated with distinct requests can share the same vehicle as long as its capacity is not exceeded. In addition, a maximum ride time is associated with each request. It corresponds to the maximum duration of the trip between the pickup point and the delivery point. This paper addresses a generalization of the DARP in which users can be transferred from one vehicle to another at intermediate points, called *transfer points* [1]. This problem is called the Dial-A-Ride Problem with Transfers (DARPT). As the DARP is a special case of the DARPT for which the set of transfer points is empty, the DARPT is NP-hard.

The DARP is a generalization of the Vehicle Routing Problem (VRP). It can be either static [2], when each request is known in advance, or dynamic, when requests can be integrated at any time [3]. Since the DARP deals with people transportation, specific requirements concerning the quality of service have to be taken into account. The objective function to be optimized is application-dependent: the number of vehicles used, total ride time, total distance, cost or service-related criteria [4]. In the remainder of this paper, we consider a static case problem in which the objective function is to minimize the total distance traveled. Quality of service is enforced by defining appropriate maximum ride times for each user.

This study is motivated by a demand-responsive transport application for people with mental disabilities traveling to and from their home to social centers, vocational rehabilitation centers or schools. These persons are able to get onto or off vehicles without physical assistance, but due to their lack of autonomy, they cannot use public transportation systems. They travel in dedicated vehicles whose use is very costly. Some organizations, responsible for the transportation of people with disabilities, wish to pool their transportation system in order to decrease costs. The objective is thus to design routes in order to service the users of several centers at a reduced cost. In order to improve the use of the vehicles, one method consists in concentrating several vehicles in predetermined locations, called transfer points, where users are transferred from one vehicle to another.

In this paper, we show that the introduction of transfer points may generate significant savings for both real-life and generated instances of the DARP. We extend an Adaptive Large Neighborhood Search (ALNS) previously designed for the Pickup and Delivery Problem with Transfers (PDPT) [5]. We also integrate a preliminary study on feasibility algorithms for the DARPT [6]. Besides, when the operators that introduce transfers are not used, the proposed ALNS competes with the state of the art algorithms for the DARP.

The remainder of this paper is organized as follows. First, we present a review of related works. Then, we formally describe the DARPT and the reasons for studying this problem. Section 4 is devoted to the description of the solution method. Section 5 focuses on the problem of checking whether a given set of routes is feasible or not. More specifically, we discuss how the feasibility can be verified within the solution method framework. Section 6 presents computational experiments, while the conclusion appears in the last section.

---

* Corresponding author. Tel.: +33 251858313; fax: +33 251858349.
E-mail addresses: olivier.peton@emn.fr, olivierpeton@yahoo.fr (O. Péton).

2. Related works

As far as we know, the only known definition of the DARPT as it is considered in this paper is given in works [1,6]. This section presents relevant studies concerning the DARP and the Pickup and Delivery Problem with Transfers (PDPT), since these problems are related to the DARPT. The PDPT, which is a special case of the DARPT without the maximum ride time constraints, has been the subject of little research. Some of these works concerning the PDPT come from Dial-A-Ride applications, where no quality of service constraints or objective is considered.

Cordeau and Laporte [2] provide a literature review about the DARP. The main heuristic methods proposed for solving the DARP are the following: Cordeau and Laporte [7] develop a Tabu Search. Parragh et al. [8] propose a Variable Neighborhood Search (VNS), whose results are competitive with the Tabu Search of Cordeau and Laporte [7]. Jain and Van Hentenryck [9] develop a constraint-based Large Neighborhood Search (LNS) to solve the DARP in limited time. Parragh and Schmid [10] integrate column generation in a Large Neighborhood Search.

As for the VRPTW, feasibility checking in local search methods has to be efficiently enforced for the DARP. The maximum ride time constraint complicates this task significantly and several algorithms have been proposed on this subject [11,12].

To the best of our knowledge, the first work considering the opportunity of transferring users in a Dial-A-Ride application is that of Stein [13]. In this study, neither the capacity of vehicles nor the time windows are taken into account. The objective is to minimize the completion time of all routes. The author proposes two heuristics for this problem. Shang and Cuff [14] present a heuristic for a PDPT in which every vertex can be used to perform a transfer. In this work, the transfer operation can be viewed as a recourse policy, as transfers are used only if the insertion of a request into the solution requires the use of an additional vehicle. Thangiah et al. [15] adapt the heuristic of Shang and Cuff to a dynamic version of the problem. Cortés and Jayakrishnan [16] develop a simulation for a demand-responsive transit system that allows a single transfer per passenger, in order to evaluate the feasibility of such a system. Mues and Pickl [17] present a heuristic column generation for a problem with a single transfer point. The algorithm is evaluated on instances with up to 70 requests. Mitrović-Minić and Laporte [18] present a local search for a PDPT with uncapacitated vehicles. Experiments are performed on instances using the Manhattan distance. These instances consider up to 100 requests. Gørtz et al. [19] propose heuristics for a PDPT with capacitated and uncapacitated vehicles. Their objective is to minimize the makespan. Petersen and Ropke [20] present an ALNS for a pickup and delivery problem with a single transfer point, concerning the transportation of flowers in Denmark. The algorithm is evaluated on real-life instances containing up to 982 requests. Qu and Bard [21] propose a GRASP coupled with an ALNS. The algorithm is evaluated on instances with up to 25 requests. Masson et al. [5] develop an ALNS for the PDPT. They present competitive results on the instances of Mitrović-Minić and Laporte and on real-life instances with up to 193 requests.

Very few exact approaches have been proposed for the PDPT. Cortés et al. [22] introduce the first mathematical formulation of it in which they use to solve instances with up to six requests using a *Branch-and-Cut* algorithm. Kerivin et al. [23] present a mathematical formulation for a version of the problem without time windows. In their formulation, the time is discretized and every vertex can be used to perform a transfer. Instances with up to 15 requests are solved using a *Branch-and-Cut* algorithm in less than 5 h. Fugenschuh [24] works on a school bus problem where the routes are already designed and need to be assigned to a minimum number of vehicles. Some passengers can be transferred from one bus to another. Nakao and Nagamochi [25] present a lower bound for the cost of the solution of the PDPT with no time window and a single transfer point. This lower bound depends on the value z (*PDP*) of the optimal solution without transfers. If $z(PDPT)$ denotes the optimal solution of the problem with transfers and $|R|$ the number of requests, then $z(PDPT) > z(PDP)/6\lceil\sqrt{|R|}\rceil + 1$. Masson et al. [26] develop a Branch-and-Cut-and-Price for a special case of the PDPT, called the Pickup and Delivery Problem with Shuttles routes, where the set of distinct delivery locations is limited. In this problem, a vehicle is not allowed to visit more than two distinct delivery locations within its route. Real-life instances are used for experimentations. Instances with up to 87 customers are solved to optimality in less than 1 h.

3. The Dial-A-Ride Problem with Transfers

In this section, we discuss the reason for considering the use of transfers in a Dial-A-Ride applications and we formally present the DARPT. The main motivation for transferring users is to achieve savings through flow consolidation at transfer points. The example depicted in Figs. 1 and 2 illustrates the savings due to transfers. Fig. 1 illustrates a solution without any transfer. Each pickup point (p_i) and the associated delivery point (d_i) are serviced by the same vehicle. Time windows are represented by the intervals above each vertex. Fig. 2 presents a solution where the requests (p_2, d_2) and (p_4, d_4) are transferred at point τ. This solution is 20% cheaper than the solution with no transfer.

3.1. Formulation of the DARPT

The DARPT can be formally described as follows. Let R be the set of requests, T the set of transfer points and K the set of



Fig. 1. A solution without transfer (DARP).

Fig. 2. A solution with transfer (DARPT).

homogeneous vehicles of capacity Q. A request $r \in R$ corresponds to the transportation of q_r passengers from a pickup point p_r to a delivery point d_r. The sets of all pickup and delivery points are denoted P and D respectively. Each vehicle performs a single route. The starting and ending depots of vehicle $k \in K$ are designated by o_k and o'_k, respectively. The set of all starting (resp ending) depot locations is denoted O (resp. O'). The duration of a route cannot exceed L. The DARPT is defined on a complete directed graph $G = (V, A)$ where $V = P \cup D \cup O \cup O' \cup T$ is the set of all vertices and $A = \{(i, j) | i, j \in V, i \neq j\}$ is the set of all arcs. A non-negative travel time $\theta_{i,j}$ and a travel cost $c_{i,j}$ are associated with each arc $(i, j) \in A$. A time window $[e_i, l_i]$ is associated with each vertex $v_i \in V$, where e_i and l_i represent the earliest and the latest times at which service can begin at vertex i. Each vertex $i \in P \cup D \cup T$ has a known service duration s_i, modeling the time needed to get users onto or off vehicle. The number of users carried simultaneously by a vehicle cannot exceed its capacity Q. A vehicle is allowed to wait at a vertex in order to service it within its time window. Each request $r \in R$ has a maximum ride time $\overline{L}_r$ which is the maximum time allowed between the end of service at p_r and the beginning of service at d_r. If two requests have a common pickup or delivery location, the corresponding vertex is duplicated. According to this modeling, pickup and delivery vertices are associated with exactly one request.

A solution of the DARPT is a set of $|K|$ routes serving all requests and such that each vehicle $k \in K$ starts at o_k and ends at o'_k. For every request $i \in R$, vertices p_i and d_i can be served by the same route, provided that p_i is served before d_i. Vertices p_i and d_i can also be served by distinct routes $k_1 \in K$ and $k_2 \in K$. In this case, k_1 and k_2 have to service a common transfer point $\tau \in T$, such that p_i is serviced before τ in k_1 and τ is serviced before d_i in k_2. Moreover the service time of τ in k_2 must occur after its service time in k_1, plus the duration of the transfer operation.

The mathematical formulation of Cortés et al. [22] for the PDPT can be extended to the DARPT by adding maximum ride time and maximum route time constraints. However, since this model contains 31 sets of constraints, we do not report it here.

3.2. Modeling transfer points

To keep track of the path followed by transferred requests, transfer points are duplicated as follows: for each request $i \in R$ transferred from route $k_1 \in K$ to route $k_2 \in K$ at transfer point $\tau \in T$, τ is duplicated in an inbound transfer vertex $t^-_{\tau,i}$ and an outbound transfer vertex $t^+_{\tau,i}$. The inbound transfer point $t^-_{\tau,i}$ represents the unloading of request i at transfer point τ, while $t^+_{\tau,i}$ models the reloading of request i at transfer point τ. This model is described in Fig. 3, where the left part illustrates the service to the physical locations and the right part shows the vertices used in the model.

4. An adaptive large neighborhood search

In this section, we describe the ALNS used to solve the DARPT. The ALNS extends the Large Neighborhood Search (LNS) introduced by Shaw [27] in a constraint programming framework to solve the Vehicle Routing Problem with Time Windows (VRPTW).

The reader interested in an extensive description of the LNS and its application to combinatorial problems is referred to the review by Pisinger and Ropke [28].

The ALNS includes an adaptive layer which enables some parameters to be automatically adjusted according to the performance of the heuristic in the last iterations. It was introduced by Pisinger and Ropke to solve a large variety of vehicle routing problems [29] including the Pickup and Delivery Problem with Time Windows [30]. It has been proven efficient for solving the PDPT [5].

4.1. Main scheme of the ALNS

The underlying principle of the ALNS is to destroy and repair a solution iteratively in order to improve it. To do so the ALNS relies on heuristic operators which aim to either destroy (removing requests from routes) or repair (reinserting requests) the solution. The general functioning of the ALNS is depicted in Algorithm 1.

Algorithm 1. ALNS

Require: *InitialSolution*
1: *BestSolution ← InitialSolution*
2: *CurrentSolution ← InitialSolution*
3: **while** the termination criterion is not satisfied **do**
4:      Selection of a *Destroy* and a *Repair* operator according to past performances
5:      *S ← CurrentSolution*
6:      *S ← Destroy(S)*
7:      *S ← Repair(S)*
8:      **if** $S \prec BestSolution$ **then**
9:          *BestSolution ← S*
10:         *CurrentSolution ← S*
11:     **else**
12:         **if** *AcceptanceCriterion(S, CurrentSolution)* **then**
13:             *CurrentSolution ← S*
14:         **end if**
15:     **end if**
16: **end while**
17: **return** *BestSolution*

The algorithm starts from an initial solution. At each iteration a pair of destroy and repair operators are selected from a pool of operators in order to create a new solution by modifying the current solution (line 4). The solution is destroyed by removing some of its requests (line 6). The repair step consists in reinserting the destroyed requests (line 7). When the resulting solution is worse than the current solution, an acceptance criterion similar to that of simulated annealing or record-to-record travel, can be used to determine whether the new solution should replace the current solution (lines 12–13). In our implementation, we use simulated annealing as an acceptance criterion. In the end, the algorithm returns the best solution encountered during the search (line 17).

The destroy and repair operators are selected with a roulette wheel selection principle: a weight is calculated for each operator and the roulette wheel selects each operator with a probability


Physical flow of a transfer operation   Corresponding graph representation

Fig. 3. Modeling of transfer operations.

that is proportional to its weight. A weight is calculated with rules similar to those described in [30]. In the beginning of the algorithm, all operators have identical weights. Then, for each segment of 100 iterations, each operator is evaluated with a score, initialized to 0 and updated as follows:

1. *New best solution*: Each time the use of an operator results in a new best known solution, its score is increased by 33.
2. *New improving solution*: Each time the use of an operator results in a new solution which improves the current one, its score is increased by 20.
3. *New accepted solution*: Each time the use of an operator results in a new solution whose cost is worse than the current solution, but accepted by the acceptation criterion, its score is updated by 15.

Operators weights are updated every 100 iterations, as an exponential smoothing of the score with a smoothing factor of 0.1.

4.2. Destroy operators

Our implementation of the ALNS uses six distinct destroy operators. These operators destroy the solution by removing requests from routes. We mainly describe the operators that are dedicated to problems with transfers. Those arising from other papers about the pickup and delivery problem are briefly enumerated and commented on.

4.2.1. Destroy operators dedicated to transfers
The objective of the following operators is to remove either transferred requests from the solution or requests that would benefit from a reinsertion with transfer. These operators were initially introduced to solve the PDPT [5].

- *Transfer point removal*: This operator simultaneously removes requests that use a given transfer point. The underlying idea is to give them a chance to be rerouted through another transfer point or with no transfer.
- *Pickup/delivery cluster removal*: This operator aims to remove a set of requests that would be efficiently routed through a common transfer point. A set of requests may benefit from using a common transfer point if their pickup (or their delivery) points form a cluster. If the pickup points of a set of requests are close enough, they can be serviced together by a common vehicle and then carried to a transfer point where they are spread among distinct vehicles.
- *History removal*: This operator is adapted from the two *history removal* heuristics proposed by Pisinger and Ropke [29]. Its objective is to remove requests that do not seem well placed in the current solution with regard to their position in the best known solutions encountered during the search. We denote $\rho(i)$ and $\sigma(i)$ as the predecessor and the successor of a vertex i within its route. For each vertex $j \in V$ and $j' \in V$, we evaluate $\xi_{j,j'}$ the number of times in the 50 best solutions that $\sigma(j) = j'$. A score ϕ_r is associated with each request $r \in R$. It is calculated as follows:
  - if request r is not transferred, $\phi_r = \xi_{\rho(p_r),p_r} + \xi_{p_r,\sigma(p_r)} + \xi_{\rho(d_r),d_r} + \xi_{d_r,\sigma(d_r)}$;
  - if request r is transferred at point $\tau \in T$, $\phi_r = \frac{1}{2}\Big(\xi_{\rho(t_{\tau,r}^-),t_{\tau,r}^-} + \xi_{t_{\tau,r}^-,\sigma(t_{\tau,r}^-)} + \xi_{\rho(t_{\tau,r}^+),t_{\tau,r}^+} + \xi_{t_{\tau,r}^+,\sigma(t_{\tau,r}^+)} + \xi_{\rho(p_r),p_r} + \xi_{p_r,\sigma(p_r)} + \xi_{\rho(d_r),d_r} + \xi_{d_r,\sigma(d_r)}\Big).$

4.2.2. General destroy operators
We implemented the following general purpose destroy operators which are common operators for the LNS and ALNS algorithms.

A good description of these heuristics can be found e.g. in [30] in the case of the Pickup and Delivery Problem with Time Windows.

- *Worst removal*: This operator iteratively removes requests that cause the biggest detour in the current solution.
- *Random removal*: This operator randomly selects the requests to be removed among the planned requests.
- *Related removal* (*or Shaw removal*): This operator aims to remove requests that are alike, so that their positions of service can somehow be interchanged.

4.3. Repair operators

The destroy operators result in a partial solution and a list of unplanned requests (also called a request bank). The purpose of repair operators is to reinsert the unplanned requests into the partial solution. We use two categories of operators: those which consider the possibility of using transfer points and those which do not resort to transfer points.

4.3.1. Repair operators with transfers

- *Best insertion with transfer*: This operator is an adaptation of the operator of Mitrović-Minić and Laporte [18] for the PDPT. For each unplanned request the best insertion without considering any transfer opportunity is first evaluated. Then for each transfer point τ and each unplanned request (p_i, d_i), we evaluate the insertion cost as follows:
  (i) the insertion cost of the pair (p_i, τ) is evaluated and considered as inserted at its best position. Then, an evaluation of the insertion cost of the pair (τ, d_i) is carried out.
  (ii) the insertion cost of the pair (τ, d_i) is evaluated and considered as inserted at its best position. Then, an evaluation of the insertion cost of the pair (p_i, τ) is carried out for every route.

  Finally, the best insertion among all evaluated insertions is performed.
- *Transfer first*: The transfer first neighborhood gives priority to the use of transfer points. As long as unplanned requests remain, best insertions with transfers are performed. If no feasible insertion with transfer can be found, the best insertion without transfer is then considered. When all requests have been inserted, a post-processing step consists in iteratively removing each transferred request and trying to reinsert it without transfer. This step aims to detect forced transfers that reduce the quality of the current solution.
- *Regret insertion with transfer*: This heuristic facilitates the insertion of requests for which using a transfer point is cheaper than insertion without transfer. For each destroyed request we compute the difference between the insertion cost of this request using a transfer point and the best insertion cost without transferring the request. The request with the largest difference is inserted first at its best position.

4.3.2. General repair operators
The following operators are those for the PDPTW, from various papers, which insert requests without transfers.

- *Best insertion*: At each iteration, the best insertion cost is computed for each unplanned request and the request with the lowest insertion cost is inserted at its best position. The heuristic stops when all requests are routed or none can be inserted anymore [30].

- *Regret heuristic*: This heuristic is based on the notion of regret, used for example by Potvin and Rousseau [31] for the Vehicle Routing Problem with Time Windows (VRPTW), and extended by Ropke and Pisinger [30]. Let U be the set of unplanned requests and, for each request $i \in U$, let Δf_i^j be the insertion cost of i in the jth best route at the best position. At each iteration, the request $i^\star$ selected for insertion at its best position is chosen such that $i^\star = \arg\max_{i \in U}(\sum_{j=2}^k \Delta f_i^j - \Delta f_i^1)$. The heuristic stops when U is empty or no request can be inserted anymore. In our implementation, we consider regret-k heuristics with values of k between 2 and 5.

5. Feasibility of a set of routes

A solution of the DARPT is feasible if it respects the capacity constraints as well as the temporal constraints of the problem. Since the way the capacity constraints are handled in the DARPT does not differ from the PDPTW, this issue is not discussed in this section. We highlight the feasibility check of a solution with regard to the temporal constraints. A solution of a Vehicle Routing Problem is feasible if and only if each route of the solution is feasible. Concerning the DARP, Cordeau and Laporte [7] proposed an algorithm to determine the feasibility of a route in $O(n^2)$, where n is the number of vertices in the route. However, in the DARPT, the routes are connected through transfer points. The need for synchronization at transfer points leads to an interdependence problem [32], so that checking the feasibility of the routes independently of each other is no longer possible.

In this section, we show that checking whether a solution of the DARPT satisfies the temporal constraints of the problem can be modeled as a Simple Temporal Problem (STP), which is a special case of the Temporal Constraint Satisfaction Problem (TCSP) [33]. Then we discuss how to solve this feasibility problem and, more specifically, how to check whether the insertion of a request in a feasible solution produces a feasible solution. Millions of feasibility checks are performed during the execution of the ALNS so this check has to be as efficient as possible. As a result, we have developed sufficient conditions and necessary conditions to accelerate the detection of feasible or infeasible solutions.

5.1. Formulation of the feasibility problem

The difficulty induced by transfers lies in the introduction of generalized precedence constraints between vertices that are serviced by different routes. The consequence is that a modification in a given route can impact the feasibility of another route. This is illustrated in Fig. 4: an insertion of a vertex before p_1 may delay the service at d_3, which can lead to a violation of the maximum ride time constraint of request 3. Therefore, a modification of a route can impact the feasibility of another route.

The set of service times at each vertex in a solution is called the schedule of the solution and denoted H. Let us denote H_i the

service time at vertex $i \in V$. The set of requests transferred at transfer point $\tau \in T$ is represented by R_τ. Let us define the route feasibility problem (*FP*), which consists in finding feasible dates H_i. (*FP*) is defined by the following:

$$H_{\sigma(i)} \geq H_i + s_i + \theta_{i,\sigma(i)}, \quad \forall i \in V \backslash O' \tag{1}$$

$$H_{t_{\tau,r}^+} \geq H_{t_{\tau,r}^-} + s_{t_{\tau,r}^-}, \quad \forall \tau \in T, \ r \in R_\tau \tag{2}$$

$$e_i \leq H_i \leq l_i, \quad \forall i \in V \tag{3}$$

$$H_{d_r} - (H_{p_r} + s_{p_r}) \leq \overline{L}_r, \quad \forall r \in R \tag{4}$$

$$H_{o_k'} - H_{o_k} \leq L, \quad \forall k \in K \tag{5}$$

$$H_i \geq 0, \quad \forall i \in V. \tag{6}$$

Constraint (1) establishes that the service time at the successors of $i \in V \backslash O'$ is greater than or equal to the service time at vertex i, plus the service duration s_i and the traveling time between these two vertices. Constraint (2) ensures the precedence between the service time at the inbound and outbound transfer vertices. Constraint (3) guarantees that vertices should be served within their time windows. Constraints (4) and (5) model the maximum ride time and maximum route time constraints respectively.

The problem (*FP*) can be modeled as a special case of the TCSP, the Simple Temporal Problem. The TCSP considers a set of variables X_i with continuous domains. Each variable represents a time moment. Every constraint l of the TCSP is represented as a set of continuous intervals $[a_l, b_l]$. Two kinds of constraints are considered: unary constraints $(a_1 \leq X_i \leq b_1) \vee \cdots \vee (a_n \leq X_i \leq b_n)$ and binary constraints $(a_1 \leq X_i - X_j \leq b_1) \vee \cdots \vee (a_n \leq X_i - X_j \leq b_n)$. In the STP, each constraint is represented by a single interval. A dummy variable X_0 is introduced to represent the beginning of the planning horizon ($X_0 = 0$). Therefore every constraint of the STP can be represented as a binary constraint. Let V' designate the set of variables and A' the set of constraints. The STP can be formulated by the following equations:

$$X_j - X_i \leq \delta_{i,j}, \quad \forall (i,j) \in A' \tag{7}$$

$$X_i \geq 0, \quad \forall i \in J \tag{8}$$

$$X_0 = 0 \tag{9}$$

Let us introduce vertex 0 representing the beginning of the planning horizon, and problem (*FP'*) defined by the following equations:

$$H_i - H_{\sigma(i)} \leq -s_i - \theta_{i,\sigma(i)}, \quad \forall i \in V \backslash O' \tag{10}$$

$$H_{t_{\tau,r}^-} - H_{t_{\tau,r}^+} \leq -s_{t_{\tau,r}^-}, \quad \forall r \in R_\tau \tag{11}$$

$$H_0 - H_i \leq -e_i, \quad \forall i \in V \tag{12}$$

$$H_i - H_0 \leq l_i, \quad \forall i \in V \tag{13}$$

$$H_{d_r} - H_{p_r} \leq \overline{L}_r - s_{p_r}, \quad \forall r \in R \tag{14}$$

$$H_{o_k'} - H_{o_k} \leq L, \quad \forall k \in K \tag{15}$$

$$H_i \geq 0, \quad \forall i \in V. \tag{16}$$

$$H_0 = 0, \quad \forall i \in V. \tag{17}$$

Constraints (10) and (11) are equivalent to constraints (1) and (2). Constraint (12) implies that the service at vertex $i \in V$ cannot start before the beginning of its time window (recall that H_0 is 0). Symmetrically, constraint (13) state that the service at vertex $i \in V$ cannot start after the end of its time window. Therefore constraints (12) and (13) are equivalent to constraint (3). Constraints


Fig. 4. Illustration of the interdependence between routes.

(14) and (15) are equivalent to constraints (4) and (5). As a matter of fact, (FP′) is a simple temporal problem which models the same feasibility problem as (FP).

5.2. Solving the feasibility problem

Evaluating route feasibility for the DARPT is equivalent to proving the consistency of the corresponding STP. The STP can be represented by a directed graph $G' = (V', A')$ – called a *distance graph* – in which each vertex represents a time variable and each arc represents a constraint between two variables. An arc from a vertex $x_i \in V'$ to a vertex $x_j \in V'$ has a weight δ_{ij}. Fig. 5 shows such a graph. For the sake of readability, only a small number of the constraints are represented.

In this figure, each kind of constraint is represented by a distinct arc format. Dotted arcs represent maximum ride time constraints. Dashed arcs are time window constraints. Plain arcs represent precedence constraints An instance of the STP is consistent if and only if its corresponding distance graph has no circuit with a negative length. Such a circuit is also called an infeasible loop [34] and corresponds to an inconsistent set of constraints.

Several algorithms have been developed to search for negative length circuits in graphs. Cherkassky et al. [35] detailed some of them, but their experimental comparisons did not lead to a clear conclusion that one is better than another. We have selected the BFCT algorithm [36], which is a variation of the Bellman–Ford algorithm. It takes graph G' as an input and indicates if the graph contains a negative length circuit or not. If no negative length circuit is found, it provides an *as late as possible* schedule. The complexity of this algorithm is $O(|V'||A'|)$. In the ALNS, BFCT is called to evaluate each insertion position for each unrouted request in routes with transfers. This operation is likely to be performed millions of times during the execution of the algorithm, which may result in a large amount of CPU time spent on this procedure.

In order to reduce the number of calls to the BFCT algorithm, we present a set of necessary and sufficient conditions. The goal of these is to help determine the feasibility of an insertion without resorting to the use of the BFCT algorithm. Necessary conditions are used to identify unfeasible insertions efficiently. Sufficient conditions establish the feasibility of a route efficiently. Eventually, if none of these conditions can prove that a solution is feasible or not, the BFCT algorithm will be used. In the remainder of this section, the *candidate vertices* designate the pair of vertices that compose the request for which insertion is evaluated.

5.2.1. Necessary conditions

In order to identify infeasible insertions in a feasible partial solution, we propose two necessary conditions. In the first one, we relax the maximum ride time and maximum route time constraints and check that the time windows constraints can be respected. In the second, we relax the precedence constraints at the transfer points, and verify that each route can meet its time constraints.

NC1: Relaxation as a PDPT: The first necessary condition (NC1) is based on the relaxation of the ride time and route duration constraints. It has been shown for the PDPT that time window violations can be identified in constant time, provided a pre-processing step of worst case complexity $O(|V_s|^2)$ is that performed after each actual insertion in a partial solution s (where V_s is the set of vertices in the partial solution s) [5]. To strengthen this necessary condition, time windows can be tightened according to the following result [33]. Let λ_{ij} be the length of the shortest path between two vertices $i \in V'$ and $j \in V'$ in the distance graph G' of a consistent instance of the STP. The set of feasible values for H_i is then $[-\lambda_{i,0}, \lambda_{0,i}]$. The insertion heuristic starts from a feasible partial solution s, which is associated with a consistent instance of the STP. In addition, if s has been proven feasible by the BFCT algorithm, $\lambda_{i,0}$ and $\lambda_{0,i}$ have been calculated for all i in V_s. It is also easy to show that inserting a vertex in a route can only lengthen the shortest paths in G (because travel times satisfy the triangular inequality). As a result, the tightened time window $[-\lambda_{i,0}, \lambda_{0,i}]$ that can be deduced from s should be satisfied after any request insertion into s. In order to be able to use these tightened time windows, we compute the $\lambda_{i,0}$ and $\lambda_{0,i}$ for all $i \in V'$ after each modification (insertion or removal of a request) of the solution. The complexity of the computation of $\lambda_{i,0}$ and $\lambda_{0,i}$ for all $i \in V'$ is $O(|V'|^2)$.

NC2: Relaxation as a DARP: The second necessary condition is based on the relaxation of constraints (2), which connect routes at transfer points. For each transfer point $\tau \in T$ and each request (p_i, d_i) transferred at τ, (p_i, d_i) is split into two requests (p_i, τ) and (τ, d_i). The maximum ride times of these requests can be computed as follows: let γ_{p_i} be the ride time minus waiting times between p_i and τ in the current solution, the maximum ride time of (τ, d_i) is set to $\overline{L}_i - \gamma_{p_i}$. Applying the same principle to the ride time γ_{d_i} between τ and d_i, request (p_i, τ) is given a maximum ride time of $\overline{L}_i - \gamma_{d_i}$. Routes for which transferred requests are split can be scheduled independently using the feasibility algorithm of Cordeau and Laporte [7] for the DARP.

Definition 1. The route feasibility problem in which all transferred requests have been split is called the "split route feasibility problem". Let us denote (SFP) the "split route feasibility problem" associated with the feasibility problem (FP).

Proposition 1. *Let (FP) be a route feasibility problem for the DARPT. If (SFP) is not consistent then (FP) is not consistent.*


Fig. 5. A STP distance graph.

Proof. Let us consider a transferred request i and its two associated split requests i' and i''. $\gamma_{i''}$ designates the actual ride time minus the waiting times between the pickup and delivery vertices of i''. If the split request i' is not ride time-consistent (does not satisfy the ride time constraints) in the schedule produced by the algorithm of Cordeau and Laporte [7], then $H_{d_i} - (H_{p_i} + s_{p_i}) > \overline{L}_{i'} + \gamma_{i''}$. Since $\overline{L}_{i'} = \overline{L}_i - \gamma_{i''}$, this is equivalent to $H_{d_i} - (H_{p_i} + s_{p_i}) > \overline{L}_i$. Hence request i is not ride time-consistent in any schedule. Similarly, we can show that if i'' is not ride time-consistent, then i cannot be ride time-consistent. Therefore if (SFP) is not consistent, (FP) is not consistent. □

Of course, time windows can be tightened in (SFP) as proposed in NC1. If no inconsistency is found, a feasible schedule for this solution may exist. In this case, sufficient feasibility conditions can be established to identify feasible solutions quickly. The complexity of checking the necessary condition NC2 is $O(n^2)$, where n designates the number of vertices in the route into which the insertion is performed.

5.2.2. Sufficient conditions

The sufficient conditions aim to identify that a given insertion is feasible. The first sufficient condition evaluates whether the insertion modifies the service time of vertices already inserted in the solution, while the second condition determines if the solution service time modifications provided by NC2 for the DARP relaxation are valid for the DARPT.

SC1: Heuristic constant time rescheduling: Let s be a consistent partial solution and let us denote $\lambda^s_{i,0}$ the length of the shortest path between i and 0 in the distance graph induced by s. According to [33], a consistent schedule H^s for s is given by $H^s_i = -\lambda^s_{i,0}$ for all vertices $i \in V_s$. Let j be an unrouted request made up of vertices p_j and d_j, and let s' be the solution resulting from an insertion of j into a partial solution s which satisfies NC1.

Proposition 2. *Let s' be a solution resulting from the insertion of request j into a partial solution s to the DAPRT. s' is consistent if*:

1. $\forall i \in V_s, H^{s'}_i = -\lambda^s_{i,0}$,
2. $H^{s'}_{p_j} \in [e_{p_j}, l_{p_j}]$,
3. $H^{s}_{d_j} \in [e_{d_j}, l_{d_j}]$,
4. j is ride time-consistent in s'.

This is a very straightforward result, but it does not show how to check that inserting a request at a given position has no impact on the current schedule. Proposition 3 presents equivalent conditions that can be checked in constant time.

Proposition 3. *s' is consistent if*:

1. $H^{s'}_{d_j} = \max(e_{d_j}, -\lambda^s_{\rho(d_j),0} + s_{\rho(d_j)} + \theta_{\rho(d_j),d_j})$,
2. $H^{s}_{p_j} = \max(e_{p_j}, -\lambda^s_{\rho(p_j),0} + s_{\rho(p_j)} + \theta_{\rho(p_j),p_j}, H^{s'}_{d_j} - \overline{L}_j - s_{p_j})$

3. $H^{s'}_{p_j} + s_{p_j} + \theta_{p_j,\sigma(p_j)} \le -\lambda^s_{\sigma(p_j),0}$,
4. $H^{s'}_{d_j} + s_{d_j} + \theta_{d_j,\sigma(d_j)} \le -\lambda^s_{\sigma(d_j),0}$.

A similar result can be established if one vertex of request j is a transfer point.

SC2: DARP solution feasibility: This second sufficient feasibility condition is based on the schedule obtained after applying the route scheduling algorithm of Cordeau and Laporte to the *split route feasibility problem* during the evaluation of NC2. By definition, if NC2 is satisfied, requests that are not transferred in this route are ride time-consistent and satisfy their time window. Split requests also satisfy their time window but should be proven ride time-consistent. In addition, at transfer points, the delivery operation should occur before the pickup. A sufficient condition is obtained if the (SFP) schedule provides consistent times for these requests.

Proposition 4. *Let s be a solution that satisfies NC2 and let H be the (SFP) schedule established by the route scheduling algorithm of Cordeau and Laporte during NC2 evaluation. For all split requests $j \in V_s$, we denote p_j, $t^-_{\tau,j}$, $t^+_{\tau,j}$, d_j the pickup, inbound, outbound and delivery vertices of j in this solution. A solution s is consistent if for all split requests $j \in s$*:

1. $H_{d_j} - (H_{p_j} + s_{p_j}) \le \overline{L}_j$,
2. $H_{t^+_{\tau,j}} \ge H_{t^-_{\tau,j}} + s_{t^-_{\tau,j}}$.

The conditions of Proposition 4 can be checked for each split request in constant time.

5.2.3. Applying necessary and sufficient conditions

The way of sequencing calls to necessary conditions NC1 and NC2 and sufficient conditions SC1 and SC2 plays an important role in the efficiency of the algorithm. Sufficient condition SC1 is used after NC1, and SC2 uses the schedule produced by the Cordeau and Laporte route scheduling algorithm during NC2 evaluation. Therefore, when both necessary conditions and sufficient conditions are used, each sufficient condition is called after its associated necessary condition. Since conditions NC1 and SC1 can be verified in constant time, it is advisable to call them before NC2 and SC2 which have a complexity of $O(n^2)$ (n is the number of vertices in the route). The BFCT algorithm is called only when all necessary conditions and none of the sufficient conditions are satisfied. The sequence of calls to necessary and sufficient conditions is depicted in Fig. 6.

6. Computational experiments

The algorithm described in this paper was developed in C++ and the experiments were run on an i3-530 computer operated by
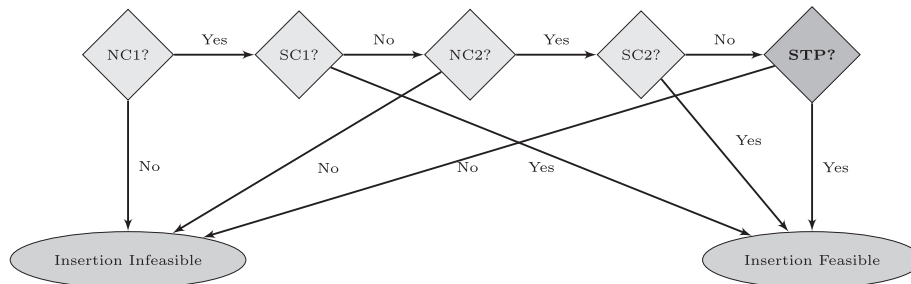

Fig. 6. Sequencing of the necessary conditions and sufficient conditions.

Ubuntu 10.04. First, we tried to evaluate the speed up due to the use of the necessary conditions and the sufficient conditions on some real-life instances. Then, we compared the solution method described in this paper to state of the art solution methods for the DARP (which is a special case of the DARPT). Finally, we evaluated the potential savings due to transfers on a set of real-life instances as well as instances from the literature.

6.1. Instances

The proposed LNS was evaluated on a set of 10 real-life instances arising from distinct sources: specialized centers for disabled children, vocational rehabilitation centers for adults and schools receiving both able-bodied and disabled children. In the first two sources, each delivery location is a common destination for dozens of requests. The instances considered correspond to a few centers which wish to pool the organization of their transport. In the following, the names of these instances start with "Centers". In the case of the transportation of disabled schoolchildren, each school is the destination for a small number of people. Thus, the instances generally include all the schools of a given geographical area, possibly several dozen. The names of these instances start with "Schools".

The instances considered include between 55 and 193 requests. Without loss of generality, we only consider outbound trips from home to the centers. Thus, the pickup points are the personal addresses of the people and the delivery points and transfer points are the centers or schools. All requests have distinct pickup points and correspond to individual demands ($q_i = 1$). The number of delivery points is comprised between 2 and 5 for specialized schools or vocational rehabilitation centers and between 13 and 33 for all other schools. Instances are named "category-$|P|-|D|$" such that the number of requests and delivery locations is reflected in their name. For example, instance Centers-193-5 is of the type "Centers" and consists of 193 requests and 5 delivery locations.

The delivery points may have different opening hours. The time window $[e_i, l_i]$ for vehicle arrival is defined such that l_i corresponds to the beginning of the daily activity. The time window opening e_i is set at 15–30 min prior to l_i, depending on the center. We assume that the vehicles start from a dummy location located at distance 0 from every pickup location. The traveling time matrices between locations are asymmetric and satisfy the triangle inequality. The maximum ride time is set at 45 min.

6.2. Efficiency of the necessary conditions and the sufficient conditions

In order to evaluate the impact of the necessary conditions and the sufficient conditions, three configurations of the method were run on the real-life instances for 1000 iterations of the ALNS. In the first configuration – called "STP" – an STP is solved with the BFCT algorithm each time one wants to evaluate whether an insertion is feasible. In the second configuration – called "NC1+NC2" – only the necessary conditions are used before resorting to the STP. In the third configuration – called "All" – all necessary conditions and the sufficient conditions are used as presented in Fig. 6. For equity reasons, all configurations evaluate exactly the same sequence of requests insertions. The results are presented in Table 1. The first column of the table indicates the instance name. The columns 2–4 report the time needed by the three configurations to perform 1000 iterations. Finally the last column reports the percentage of feasible insertions among the evaluated insertions. Note that some trivially infeasible insertions are not evaluated (e.g. if $e_i + s_i + \theta_{i,j} > l_j$ the insertion of vertex j in any position located after vertex i in a route is not evaluated). Center instances have much fewer feasible insertions because, on average, people

Table 1
Impact of the necessary conditions and sufficient conditions on the solving time.

Instances	Computation time (in s)			% of feasible insertions
	STP	NC1 + NC2	All	
Center-81-2	429	106	**83**	1.3
Center-84-2	142	62	**58**	2.8
Center-87-2	320	132	**98**	2.4
Center-109-2	1981	247	**187**	1.1
Center-193-5	15 249	1028	**900**	1.2
School-55-16	**51**	52	52	53.0
School-66-13	61	61	**59**	39.8
School-84-21	127	120	**118**	31.9
School-84-33	201	165	**162**	23.1
School-106-24	232	195	**189**	20.7
Center	18 121	1575	**1326**	1.3
School	672	593	**580**	25.7
All	18 793	2168	**1906**	2.0

Table 2
Percentage of failure of the necessary conditions and sufficient conditions.

Instances	NC1 (%)	NC2 (%)	SC1 (%)	SC2 (%)
Center-81-2	17.2	2.1	95.6	61.1
Center-84-2	41.1	1.2	95.9	42.0
Center-87-2	21.3	2.6	97.1	57.5
Center-109-2	12.9	1.3	94.5	52.4
Center-193-5	6.8	0.5	96.0	46.9
School-55-16	41.5	0.8	94.1	24.4
School-66-13	35.1	2.4	92.1	28.9
School-84-21	21.7	1.1	94.3	30.6
School-84-33	13.7	1.0	96.4	33.8
School-106-24	19.5	0.7	94.9	32.8

live further from their delivery point and the maximum ride time constraint is more binding.

The speed-up factor due to the necessary conditions and sufficient conditions varies a great deal from one instance category to another. For the *School* instances, where the percentage of feasible insertions is very high, the speed-up is almost negligible. However, for the biggest *Center* instances, the speed-up factor can reach 10 or 15. For these *Center* instances, most of the speed-up is provided by the necessary conditions.

Table 2 reports the failure rates of the necessary conditions and the sufficient conditions. It is considered that a necessary condition fails when it does not detect infeasible insertions. Symmetrically, a sufficient condition fails when it does not detect a feasible insertion. The first column is the instance name. The next columns present the failure percentage of each category of conditions.

In the worst case, the DARP relaxation (NC2) has a failure rate of 2.6%, while the PDPT relaxation (NC1) has a failure rate of 41.5%. The average failure rate of NC1 is 23.1%. However, the sufficient conditions do not exhibit such low failure rates. In the best case, the failure rate of SC1 is 92.1%, while failure rates of SC2 range from 24.4% to 61.1%. These results are not very surprising since sufficient conditions only perform well for small modifications of the current solution. On the other hand, the necessary conditions only relax a few constraints, which explains why they are able to detect most of the infeasibilities. For the instances with a very low percentage of feasible insertions, improving the sufficient conditions would have only a marginal effect. However, for instances such as the *School* instances, where the percentage of feasible

insertions ranges from 20% to 50%, improving the sufficient conditions would improve the method.

6.3. Evaluation on instances of the Dial-A-Ride Problem

There are no available benchmark instances for the DARPT. The closest optimization problem is the DARP, which is a special case of the DARPT with $T = \varnothing$. Thus, we solved the instances of Cordeau and Laporte for the DARP [7] with our implementation of the ALNS and compared its performance with algorithms solving the DARP.

6.3.1. Evaluation after 5 min of computation
First, we present the results obtained by our ALNS after 5 min of computation and compare them with the LNS-FFPA algorithm of

Jain and Van Hentenryck [9] and the Hybrid LNS of Parragh and Schmid [10]. Each instance was run five times. The results are reported in Table 3. Columns 1–3 represent the name of the instance, the number of vehicles and the number of requests, respectively. Columns 4 and 5 are the average results and the relative difference between the best known results and the average solution of the algorithm LNS-FFPA algorithm. Columns 6 and 7 report the same information for the Hybrid LNS. Finally, these results are reported for our implementation of the ALNS in columns 8 and 9.

The ALNS finds better average results in 13 out of 20 instances. The relative gap between the best known solutions and the average cost found by the ALNS is 3.74% on average. The gap with the best solutions found by the ALNS is even below 2% on average.

6.3.2. Evaluation after 50 000 iterations of the ALNS
In this experiment, the number of iterations performed by the ALNS was set at 50 000. We compared the results of the ALNS with those of the VNS of Parragh et al. [8] and of the H-LNS of Parragh and Schmid [10]. Table 4 presents the results. Columns 2–4 report the solution with the minimum cost, average cost and average run time (in minutes) of the VNS of Parragh et al. [8]. Columns 5–7 report the same information for the Hybrid LNS [10]. Columns 8–10 concern our ALNS. Column 11 reports the value of the best known solution for each instance.

On average, ALNS is competitive with the state of the art methods. The solutions found by the ALNS are, on average, around 1% worse than the best known solutions.

6.4. Savings due to transfers

This section evaluates the savings due to transfers. Since optimal solutions to instances of the DARPT of realistic size are unknown, a comparison of optimal solutions with or without transfers is impossible. Thus, we used two approaches to evaluate the savings due to transfers. First, we compared heuristic solutions of the DARPT to known optimal solutions of the DARP. This yielded lower bounds for the savings due to transfers. Then, we compared heuristic solutions of the DARP and the DARPT on real-life

Table 3
Results with limited running time.

Instance			LNS-FFPA [9]		H-LNS [10]		ALNS	
Name	$|K|$	$|R|$	Avg.	Δ_{BKS} (%)	Avg.	Δ_{BKS} (%)	Avg.	Δ_{BKS} (%)
R1a	3	24	190.77	0.39	**190.02**	0.00	**190.02**	0.00
R2a	5	48	304.45	1.03	303.63	0.76	**301.98**	0.21
R3a	7	72	547.15	2.85	536.77	0.90	**536.74**	0.89
R4a	9	96	595.05	4.35	**581.41**	1.96	588.37	3.18
R5a	11	120	662.56	5.56	**643.72**	2.56	655.02	4.36
R6a	13	144	832.74	6.05	837.01	6.59	**824.21**	4.96
R7a	4	36	292.86	0.39	291.93	0.08	**291.71**	0.00
R8a	6	72	505.15	3.23	**498.44**	1.86	500.30	2.24
R9a	8	108	711.6	7.84	697.22	5.66	**691.45**	4.79
R10a	10	144	911.18	6.55	908.12	6.19	**903.35**	5.64
R1b	3	24	**164.46**	0.00	**164.46**	0.00	**164.46**	0.00
R2b	5	48	301.67	2.03	**297.99**	0.79	300.58	1.66
R3b	7	72	504.69	4.10	**495.02**	2.10	499.26	2.98
R4b	9	96	566.48	7.02	**540.76**	2.16	554.96	4.84
R5b	11	120	610.33	5.60	607.17	5.05	**604.28**	4.55
R6b	13	144	785.13	6.43	**772.21**	4.68	780.71	5.83
R7b	4	36	248.31	0.04	**248.21**	0.00	**248.21**	0.00
R8b	6	72	477.75	3.55	473.19	2.56	**472.35**	2.38
R9b	8	108	633.51	6.75	620.39	4.54	**617.22**	4.00
R10b	10	144	857.95	8.16	874.36	10.23	**844.48**	6.46
Avg.			535.19	5.05	529.10	3.86	528.48	3.74

Table 4
Results with a limit on the number of iterations.

Instances	VNS [8]			H-LNS [10]			ALNS			BKS [10]
	Min.	Avg.	CPU	Min.	Avg.	CPU	Min.	Avg.	CPU	
R1a	**190.02**	**190.02**	5.18	**190.02**	**190.02**	0.54	**190.02**	**190.02**	2.09	190.02
R2a	**301.34**	301.78	12.06	**301.34**	302.53	2.76	**301.34**	**301.72**	6.06	301.34
R3a	533.01	536.08	19.31	535.28	538.21	5.11	**532.00**	**533.55**	13.17	532
R4a	573.92	578.21	24.13	**571.09**	**576.26**	16.29	573.47	579.54	30.47	570.25
R5a	636.77	637.72	87.69	629.52	637.59	26.70	**631.40**	**636.14**	44.00	627.68
R6a	802.12	809.27	94.10	788.88	800.35	48.48	791.35	**797.92**	64.96	785.26
R7a	**291.71**	294.26	5.09	**291.71**	292.56	1.02	**291.71**	292.56	3.45	291.71
R8a	490.58	495.7	44.29	491.93	**495.2**	5.92	**491.02**	496.01	14.76	489.33
R9a	666.96	673.18	103.97	**661.47**	676.09	24.89	663.52	**669.42**	34.43	659.85
R10a	866.97	873.15	211.29	872.31	878.93	47.17	**862.25**	**866.24**	67.75	855.15
R1b	164.46	164.46	6.24	164.46	166.06	0.61	164.46	165.52	2.19	164.46
R2b	296.65	299.19	15.71	**295.96**	298.88	3.00	**295.96**	**296.68**	9.17	295.66
R3b	490.16	494.23	28.83	**484.83**	491.29	8.19	486.99	491.83	23.82	484.83
R4b	533.15	540.5	72.66	534.84	541.19	22.58	**532.91**	**535.59**	51.26	529.33
R5b	583.12	588.48	131.85	587.67	590.22	44.09	**578.41**	**584.99**	91.54	577.98
R6b	742.28	751.55	227.25	738.01	743.64	71.50	741.70	748.75	128.36	737.69
R7b	248.21	248.21	8.77	248.21	248.21	1.30	248.21	248.21	5.06	248.21
R8b	462.50	468.97	39.10	463.67	470.25	9.54	**461.39**	**465.34**	23.63	461.39
R9b	600.18	607.94	75.70	**593.49**	606.25	27.49	600.19	604.66	68.13	593.49
R10b	796.90	805.93	250.20	804.22	812.81	69.57	**795.64**	**799.64**	118.18	793.21
Avg.	513.55	517.94	73.17	512.45	517.83	21.84	511.70	515.22	40.12	509.44

Table 5
Savings due to transfers on instances adapted from Ropke et al. [37].

Inst.	B&C $T=\varnothing$	ALNS $T=\{Depot\}$				ALNS $T=N$			
	Opt.	Avg.	Min.	Gap no tr. (%)	CPU (s)	Avg.	Min.	Gap no tr. (%)	CPU (s)
a2-16	294.25	294.25	294.25	0.00	70	293.70	293.70	−0.19	251
a2-20	344.83	344.83	344.83	0.00	106	340.96	340.96	−1.14	644
a2-24	431.12	412.80	412.80	−4.44	160	410.00	410.00	−5.15	1013
a3-24	344.83	341.88	341.88	−0.86	149	338.71	338.32	−1.92	938
a3-30	494.85	485.61	485.59	−1.91	292	476.42	476.09	−3.94	2104
a3-36	583.19	548.43	548.31	−6.36	497	531.90	531.49	−9.73	4466
a4-40	557.69	557.94	557.94	0.04	527	544.86	543.81	−2.55	4989
a4-48	668.62	666.66	665.87	−0.41	787	639.33	638.06	−4.79	13 675
a5-40	498.41	496.54	496.54	−0.38	462	490.47	490.46	−1.62	6122
a5-50	686.62	671.00	669.43	−2.57	1106	649.97	646.94	−6.13	17 527
a5-60	808.42	803.40	802.86	−0.69	1306	779.88	778.17	−3.89	33 126
a6-48	604.12	608.85	608.22	0.67	804	595.49	592.89	−1.89	11 586
a6-60	819.60	800.08	796.74	−2.87	1454	779.10	775.23	−5.72	32 722
a6-72	916.05	919.12	914.71	−0.15	2426	908.28	904.85	−1.24	t.o.
a7-56	724.04	717.60	715.38	−1.21	961	697.14	690.85	−4.80	21 731
a7-70	889.12	902.43	897.96	0.98	3246	876.38	872.18	−1.94	t.o.
a7-84	1033.37	1040.12	1035.56	0.21	3681	1028.87	1020.90	−1.22	t.o.
a8-64	747.46	737.06	733.44	−1.91	1625	720.55	717.84	−4.13	t.o.
a8-80	945.73	939.50	936.17	−1.02	2968	916.60	909.75	−3.95	t.o.
a8-96	1232.61	1249.32	1238.84	0.50	6517	1210.72	1183.73	−4.13	t.o.
Avg.	681.25	676.87	674.87	−1.12	1457	661.47	657.81	−3.50	18 345

instances. This estimated the savings achievable in practice in the context of demand-responsive transport.

6.4.1. Lower bounds for the savings due to transfers

In order to compute lower bounds for the savings due to transfer points, we compared heuristic solutions to the DARPT with optimal solutions to the DARP computed by a Branch-and-Cut algorithm [37]. For each DARP instance we created two DARPT instances. The first one considers one unique transfer point located at a depot in the center of the area. The second one considers every vertex as a potential transfer point.

Table 5 reports the results of five runs of the ALNS limited to 25 000 iterations or 10 h. Column 1 is the instance name. Column 2 reports the optimal solution of the DARP. Columns 3–6 report the average solution, the best solution, the savings due to transfers and the run time of the ALNS with a single transfer point (in seconds) respectively. Columns 7 to 10 report similar information in the case of multiple transfer points.

Savings due to transfers vary a great deal from one instance to another. Moreover, the location and number of the transfer points seem to have a non-negligible impact on savings. When only the depot can be used to perform a transfer, the savings are more than 1.5% in only 6 instances. When multiple transfer points are authorized, the savings are more than 1.5% in 14 instances. On average the savings are 3 times bigger in the latter case. However, solving times tend to be very long. The limit of 10 h is exceeded – written "t.o." in the table – when the potential number of transfer points is larger than 140 requests (70 pickup points + 70 delivery points), when all points can be used to perform a transfer. Note that in some cases, the ALNS for the DARPT cannot find better solutions than the optimal solution to the DARP. These results can be explained by the structure of the instance of Cordeau and Laporte, which is not really adapted to transfers.

6.4.2. Savings on real-life instances

In order to compute savings due to transfers on real-life instances, we compared heuristic solutions of the DARP and the DARPT, both computed with the ALNS. Each instance was solved

Table 6
Evaluation of savings due to transfer on real cases.

Instances	$|T|$	Avg. cost	Min. cost	CPU (s)	Gap tr. (%)
Center-81-2	–	499.11	496.97	1265	−4.92
	2	477.21	472.52	15 539	
Center-84-2	–	939.24	934.88	1088	−2.17
	2	919.84	914.60	7725	
Center-87-2	–	500.72	500.72	1572	−3.47
	2	488.09	483.37	16 777	
Center-109-2	–	740.40	740.08	1917	−7.19
	2	690.89	686.90	23 186	
Center-193-5	–	1347.60	1342.38	6168	−8.24
	5	1242.62	1231.82	t.o.	
School-55-16	–	672.03	672.03	480	−0.97
	16	665.58	665.53	3044	
School-66-13	–	859.52	859.52	662	−2.66
	13	836.62	836.62	4873	
School-84-21	–	1142.52	1142.52	1008	−3.25
	21	1109.16	1105.40	12 914	
School-84-33	–	1269.67	1269.67	991	−6.50
	33	1188.47	1187.08	24 270	
School-106-24	–	995.90	995.58	1475	−3.59
	24	961.55	959.82	29 721	

5 times by the ALNS with a limit of 25 000 iterations or 10 h. Table 6 reports the results. Column 1 is the instance name. Column 2 indicates the number of transfer points (– means a DARP instance). Columns 3 and 4 report the average and best solution over the 5 runs. Column 5 reports the average running time (in seconds). Finally column 6 reports the gap between the solutions with and without transfers.

The savings due to transfers range from 2.17% to 8.24% for the *Center* instances and from 0.97% to 6.50% for the *School* instances. The counterpart is a large increase in the computing time. In order to evaluate whether the maximum ride time constraint has an impact on these savings, we also solved the instances without these constraints – in this case, the problem is a Pickup and Delivery Problem with Transfers (PDPT). The results are reported in Table 7. The columns have the same meaning as columns 1–4 and 6 of Table 6.

Table 7
Solution without maximum ride time constraints.

Instances	$|T|$	Avg. cost	Min. cost	Gap tr. (%)
Center-81-2	–	453.09	452.72	−3.23
	2	438.64	438.08	
Center-84-2	–	782.04	782.03	−2.54
	2	769.83	762.15	
Center-87-2	–	478.06	477.25	−2.96
	2	463.92	463.12	
Center-109-2	–	650.20	647.88	−6.33
	2	610.74	606.85	
Center-193-5	–	1253.63	1250.30	−9.23
	5	1152.44	1134.90	
School-55-16	–	654.72	654.72	−0.75
	16	649.78	649.78	
School-66-13	–	824.82	824.82	−2.98
	13	800.51	800.27	
School-84-21	–	1121.37	1121.37	−6.17
	21	1055.67	1052.17	
School-84-33	–	1239.65	1239.32	−6.23
	33	1164.48	1162.08	
School-106-24	–	978.64	978.15	−4.12
	24	940.02	937.85	

The cost of solutions is higher when maximum ride time constraints are considered, which is understandable since the problem is more constrained. However, the savings due to transfers are rather similar with or without maximum ride time constraints. The only exception is the instance School-84-21 for which the savings due to the use of transfers is 6.17% in the case of the PDPT and 3.25% in the case of the DARPT. Therefore it seems that on those real-life instances, the maximum ride time constraint does not have a large impact on the savings due to transfers.

6.5. Performance of the destroy and repair operators

In order to assess the performance of the destroy and repair operators, we have conducted experiments on a subset of 10 representative instances: instances a2-20, a3-18, a3-30, a4-16 and a4-32 by Cordeau and Laporte, and real-life instances Center-81-2, Center-84-2, School-55-16, School-66-13 and School-84-21.

In the ALNS algorithm described in Section 4.1, the score of an operator is increased on three events: the operator provides a new best known solution, a new improving solution, or a new accepted solution. For each event, we measure the individual contribution of each operator as the following ratio: the number of times the event arises from the operator over the total number of occurrences of the event.

These results are summarized in Table 8 for destroy operators and in Table 9 for repair operators. In each table, the first column reports the name of the operators. Columns 2 to 4 report the minimal, average and maximal percentage of new best solutions found by each operator. Columns 5 to 7 report the same information concerning improvements over the current solution. Finally, columns 8 to 10 report this information for new deteriorating solutions accepted by the acceptation criterion. Note that the results in column 3, 6 and 9 sum up to 100% for each category of operators: destroy operators and repairs operators. In Table 9, the best insertion and the regret insertion results are aggregated in the last line.

Regarding destroy operators (Table 8), none clearly dominates the others. The related removal operator is responsible, on average, for the identification of 24% of the new best solutions. But on some instances, this operator is not involved in the

Table 8
Performance of the destroy operators during the execution of the ALNS algorithm.

Destroy operators	New BKS			New improving			New accepted		
	Min. (%)	Avg. (%)	Max. (%)	Min. (%)	Avg. (%)	Max. (%)	Min. (%)	Avg. (%)	Max. (%)
Transfer point removal	0	11	30	10	15	20	7	15	22
Pickup/delivery cluster removal	0	15	29	5	17	23	5	17	22
History removal	0	12	20	4	12	19	4	12	18
Worst removal	0	19	40	0	15	26	0	14	23
Random removal	10	19	29	13	20	33	13	19	27
Related removal	0	24	50	15	21	41	15	23	41

Table 9
Performance of the repair operators during the execution of the ALNS algorithm.

Repair operators	New BKS			New improving			New accepted		
	Min. (%)	Avg. (%)	Max. (%)	Min. (%)	Avg. (%)	Max. (%)	Min. (%)	Avg. (%)	Max. (%)
Best insertion with transfer	17	29	50	23	28	41	22	34	43
Transfer first	36	51	83	36	53	61	36	45	53
Regret insertion with transfer	0	13	31	1	11	23	2	12	23
Best insertion + regret	0	7	24	0	8	16	0	9	17

identification of a single new best solution. Clearly, the efficiency of each destroy operator depends on the instance.

On the contrary, the results in Table 9 clearly show the need for repair operators suited for transfers. On average, the Transfer First operator is involved in half of the occurrences of the three events. The regret insertion with transfer seems to be less efficient, but still can identify up to 31% of the new best solutions on some instances. In conclusion, the contribution of repair operators also varies a lot from one instance to another, but all operators that introduce transfers are meaningful. Even though the best insertion and the regret insertion seem less efficient, they are less time consuming than the others and they must be kept to remove some transfers.

7. Conclusion

In this paper, we developed an Adaptive Large Neighborhood Search (ALNS) to solve the Dial-A-Ride Problem with Transfers. The algorithm strongly relies on previous work carried out on the Pickup and Delivery Problem with Transfers [5] and competes with state-of-the art methods for the DARP. We show that the feasibility of the insertion of a request into a feasible solution can be efficiently checked. We provide lower bounds for the savings due to transfers and apply the ALNS to real-life instances for which we evaluate the practical savings due to transfers. The introduction of transfer points can lead to non-negligible savings (up to 8.25%). However, it does not seem straightforward to characterize instances that may benefit most from the introduction of transfers.

This work opens up some new perspectives. In this study, the quality of service provided to the users has been modeled only by maximum ride time constraints. However, users generally consider transfers as a deterioration in their quality of service. As a

matter of fact, one could consider to restricting the number of passengers transferred, the maximum waiting time of an user at a transfer point or the overall waiting time of the users at transfer points.

Another perspective concerns the interconnection between routes. When the number of interconnected routes grows, a single unexpected event in one route can propagate to the whole network. Therefore, building robust solutions to the DARPT seems to be a challenging line of research.

References

[1] Masson R, Lehuédé F, Péton O. A tabu search algorithm for the Dial-A-Ride Problem with Transfers. In: Proceedings of the international conference on industrial engineering and systems management. Metz, France; 2011. p. 1224–32.

[2] Cordeau J-F, Laporte G. The Dial-A-Ride Problem: models and algorithms. Annals of Operations Research 2007;153(1):29–46.

[3] Berbeglia G, Cordeau J-F, Laporte G. Dynamic pickup and delivery problems. European Journal of Operational Research 2010;202(1):8–15.

[4] Paquette J, Cordeau J-F, Laporte G. Quality of service in Dial-A-Ride operations. Computers & Industrial Engineering 2008;56:1721–34.

[5] Masson R, Lehuédé F, Péton O. An adaptive large neighborhood search for the pickup and delivery problem with transfers. Transportation Science, articles in advance. http://dx.doi.org/10.1287/trsc.1120.0432.

[6] Masson R, Lehuédé F, Péton O. Simple temporal problems in route scheduling for the Dial-A-Ride Problem with Transfers. In: Beldiceanu N, Jussien N, Pinson E, editors. Integration of AI and OR techniques in constraint programming for combinatorial optimization problems – 9th international conference, CPAIOR 2012, Nantes, France, May 28–June 1, 2012. Proceedings. Lecture notes in computer science, vol. 7298. Springer, 2012. p. 275–91.

[7] Cordeau J-F, Laporte G. A tabu search heuristic for the static multi-vehicle Dial-A-Ride Problem. Transportation Research Part B: Methodological 2003;37(6):579–94.

[8] Parragh SN, Doerner KF, Hartl RF. Variable neighborhood search for the Dial-A-Ride Problem. Computers & Operations Research 2010;37:1129–38.

[9] Jain S, Van Hentenryck P. Large neighborhood search for Dial-A-Ride Problems. In: Lee J, editor. Principles and practice of constraint programming – CP 2011. Lecture notes in computer science, vol. 6876. Berlin, Heidelberg: Springer-Verlag; 2011. p. 400–13.

[10] Parragh S, Schmid V. Hybrid column generation and large neighborhood search for the Dial-A-Ride Problem. Computers & Operations Research 2013;40(1):490–7.

[11] Firat M, Woeginger GJ. Analysis of the Dial-A-Ride Problem of Hunsaker and Savelsbergh. Operations Research Letters 2011;39(1):32–5.

[12] Hunsaker B, Savelsbergh M. Efficient feasibility testing for Dial-A-Ride Problems. Operations Research Letters 2002;30:169–73.

[13] Stein D. Scheduling Dial-A-Ride Transportation systems. Transportation Science 1978;12:232–49.

[14] Shang JS, Cuff CK. Multicriteria pickup and delivery problem with transfer opportunity. Computers & Industrial Engineering 1996;30(4):631–45.

[15] Thangiah S, Fergany A, Awam S. Real-time split-delivery pickup and delivery time window problems with transfers. Central European Journal of Operations Research 2007;15:329–49.

[16] Cortés CE, Jayakrishnan R. Design and operational concepts of high-coverage point-to-point transit system. Transportation Research Record 2002;1783:178–87.

[17] Mues C, Pickl S. Transshipment and time windows in vehicle routing. In: ISPAN'05: proceedings of the 8th international symposium on parallel architectures, algorithms and networks. Washington, DC, USA: IEEE Computer Society; 2005. p. 113–9.

[18] Mitrović-Minić S, Laporte G. The pickup and delivery problem with time windows and transshipment. INFOR 2006;44(3):217–28.

[19] Gørtz IL, Nagarajan V, Ravi R. Minimum makespan multi-vehicle Dial-A-Ride. In: Proceedings of the 17th ESA, Copenhagen, Denmark. Lecture notes in Compoter Science, vol. 5757, 2009. p. 540–52.

[20] Petersen HL, Ropke S. The pickup and delivery problem with cross-docking opportunity. In: Proceedings of the second international conference on computational logistics, ICCL'11. Berlin, Heidelberg: Springer-Verlag; 2011. p. 101–13.

[21] Qu Y, Bard JF. A GRASP with adaptive large neighborhood search for pickup and delivery problems with transshipment. Computers & Operations Research 2012;39(10):2439–56.

[22] Cortés CE, Matamala M, Contardo C. The pickup and delivery problem with transfers: formulation and a branch-and-cut solution method. European Journal of Operational Research 2010;200(3):711–24.

[23] Kerivin HLM, Lacroix M, Mahjoub AR, Quilliot A. The splittable pickup and delivery problem with reloads. European Journal of Industrial Engineering 2008;2(2):112–33.

[24] Fugenschuh A. Solving a school bus scheduling problem with integer programming. European Journal of Operational Research 2009;193(3):867–84.

[25] Nakao Y, Nagamochi H. Worst case analysis for pickup and delivery problems with transfer. IEICE Transactions on Fundamentals of Electronics. Communications and Computer Sciences 2010;E91-A(9):2328–34.

[26] Masson R, Ropke S, Lehuédé F, Péton O. A branch-and-cut-and-price for the pickup and delivery problem with shuttles routes. Technical report 12/3/AUTO, Ecole des Mines de Nantes, Nantes, France, 2012.

[27] Shaw P. Using constraint programming and local search methods to solve vehicle routing problems. In: Proceedings of the 4th international conference on principles and practice of constraint programming, 1998. p. 417–31.

[28] Pisinger D, Ropke S. Large neighborhood search. In: Gendreau M, Potvin J-Y, editors. Handbook of metaheuristics, international series in operations research & management science, vol. 146. USA: Springer; 2010. p. 399–419.

[29] Pisinger D, Ropke S. A general heuristic for vehicle routing problems. Computers & Operations Research 2007;34(8):2403–35.

[30] Ropke S, Pisinger D. An adaptive large neighborhood search heuristic for the pickup and delivery problem with time windows. Transportation Science 2006;40:455–72.

[31] Potvin J-Y, Rousseau J-M. A parallel route building algorithm for the vehicle routing and scheduling problem with time windows. European Journal of Operational Research 1993;66(3):331–40.

[32] Drexl M. Synchronization in vehicle routing—a survey of VRPs with multiple synchronization constraints. Transportation Science 2012;46(3):297–316.

[33] Dechter R, Meiri I, Pearl J. Temporal constraint networks. Artificial Intelligence 1991;49(1–3):61–95.

[34] Shostak R. Deciding linear inequalities by computing loop residues. Journal of the Association for Computing Machinery 1981;28:769–79.

[35] Cherkassky BV, Georgiadis L, Goldberg AV, Tarjan RE, Werneck RF. Shortest-path feasibility algorithms: an experimental evaluation. Journal of Experimental Algorithmics 2009;14:2.7–37.

[36] Tarjan RE. Shortest paths. Technical Report, AT&T Bell Laboratories, Murray Hill, NJ; 1981.

[37] Ropke S, Cordeau J-F, Laporte G. Models and branch-and-cut algorithms for pickup and delivery problems with time windows. Networks 2007;49(4):258–272.